

# What Primnox is

A plain-English introduction. No technical background assumed.

Primnox is an AI assistant that runs **on your own computer** rather than on someone else's website. Your conversations, your files and the things it remembers about you are stored on your machine, in a file you own.

It can use an AI model over the internet if you give it a key, or a model running on your own computer, in which case nothing you type leaves the building at all. That choice is yours and you can change it at any time.

## Why people keep saying V1 and V2

There are two versions, and they exist at the same time on purpose.

**V1** is the version that works today. It has the most features.

**V2** is a rebuild of the machinery underneath. Think of a car: V1 is the car you are driving; V2 replaced the engine and the frame while V1 kept driving. They run side by side, on different doors, with separate files — so work on V2 cannot break the version you actually use.

The rebuild was not about adding features. It was about one thing: in V1, when you sent a message, the app started working and that piece of work had no name. Anything without a name cannot be pointed at — so it cannot be stopped, checked on, or picked back up. In V2 every request is a numbered thing that exists in a list.

That sounds abstract. Here is what it actually changed:

| Before (V1)                                                                                                  | After (V2)                                    |
|--------------------------------------------------------------------------------------------------------------|-----------------------------------------------|
| Once it started replying, you could not stop it. There was no stop button because there was nothing to stop. | A stop button that works.                     |
| Refresh the window mid-answer and the answer was gone.                                                       | It picks up exactly where it was.             |
| When the AI service failed, the error appeared as a message from Primnox — as if the assistant had said it.  | Errors look like errors, with a retry button. |
| Sending a second message made it wait for the first to finish completely.                                    | Requests are tracked separately.              |

Information spread across three separate files.

One file. Everything saved together or not at all.

## What it can do today

- **Hold a conversation**, and use tools while doing it — run code safely in a sealed area, read files you give it, search your documents.
- **Make real files**. PowerPoint decks, PDFs, Word documents and spreadsheets, which you can preview inside the app and download.
- **Remember things about you** between conversations, and show you a list of everything it believes, so you can delete any of it.
- **Read a folder** of documents or code and answer questions about it with references back to the exact file, rather than guessing.
- **Ask you a question** when it is genuinely unsure, instead of picking an answer and sounding confident about it.
- **Ask permission** before doing anything that touches your computer, and wait for you.

Roughly 535 automatic checks run against V2 every time something changes. They take about a minute and they are the reason a fix in one place is not quietly breaking something in another.

## What is NOT there yet

This is the section to read carefully, because it is the one you would otherwise assume was handled.

### Privacy filtering is not built in V2

**V1 has a privacy filter.** Before anything is sent to an AI service on the internet, it looks through the text for personal details — names, email addresses, phone numbers, account numbers — and replaces them, then puts the real ones back in the answer afterwards.

**V2 has the doorway for that filter, but the filter itself is not installed.** Everything is routed through a single point so it can be added in one place later — but at that point today, the text is passed along exactly as you typed it.

So, plainly: **in V2, if you are using an AI service over the internet, what you type is sent as written.** If you are using a model running on your own computer, none of it leaves your machine and there is nothing to filter.

This is deliberate rather than forgotten. The empty filter was left obviously empty instead of half-written, so that nothing in the app could claim a protection it was not actually providing. It is planned work, not a decision that it does not matter.

## Nothing checks who is asking

Primnox only listens to your own computer — nobody on the internet can reach it. But it does not yet ask for any kind of password, so any other program already running on your machine could talk to it and read your conversations. On a personal computer that you control, this is a small risk. It is still a real one, and it is known.

## Features V1 has and V2 does not

Calendar, notes, meetings, research and the guided setup all exist in V1 and have not been rebuilt yet. V2 does fewer things on purpose — the rebuild was of the foundation, and the rooms get built back on top of it.

# Where it is going

The plan after V2, in the same plain English.

Two more versions are planned. Each one is a single idea, not a pile of features.

## V3 — an assistant that keeps working

Today you open Primnox and ask it something. The goal of V3 is that you **give it a job and walk away**.

Instead of a chat that scrolls past, a job like *“build my website”* becomes a **Mission** — a thing with a name and a progress list you can watch. Planning: done. Design: done. Building the pages: happening now. Testing: next.

It keeps going if you close the lid, restart, or come back tomorrow. It stops and asks you before anything it should not decide alone.

### Also planned for V3:

- **Several specialists instead of one.** A planner, a builder, a designer, a tester — and you can see which one is doing what, rather than one silent box.
- **It knows what you are looking at.** If you turn that on, saying “fix this” works, because it already knows which window and which file you mean. No copying and pasting.
- **Talking to it out loud**, including interrupting it mid-sentence the way you would a person.
- **Reading your screen properly** — recognising a dialog box or an error message rather than just seeing a picture.
- **A map of what it knows about you**, that you can walk through: people, projects, documents, decisions — and for each one, where it learned that.

— **Matching how you work.** Short answers while you are heads down, fuller explanations while you are learning.

## V4 — the same assistant on every device

One brain, every device. Your computer, your laptop, your phone, a browser tab — the same assistant, the same memory, the same jobs in progress.

— **Start on one device, finish on another.** Begin something at your desk, check it from your phone on the bus, finish on a laptop. It does not start over.

— **Your devices find each other directly**, over your own network, without going through anybody's cloud.

— **Your phone becomes the remote control** — approve things, watch progress, get alerted when it needs you.

— **Everything still works with no internet.** Syncing between devices is the only part that needs a connection.

— **Anything that does travel is encrypted end to end**, so even if it passes through a server, that server cannot read it.

V4 is also where the privacy filter matters most — once information moves between devices, the protection described above stops being optional. It is planned before then, in V2's own follow-up.

## The short version

| Version    | In one line                                                                                                       |
|------------|-------------------------------------------------------------------------------------------------------------------|
| V1 — today | Does the most. Built in a way that cannot be stopped, resumed or checked on.                                      |
| V2 — now   | Does less, but every request is a real, named thing. This is the foundation. Privacy filtering not yet installed. |
| V3 — next  | You assign missions instead of asking questions. Voice, screen awareness, a team of specialists.                  |
| V4 — after | The same assistant follows you across every device, and works offline.                                            |
| V5 — later | Anticipating what you need before you ask.                                                                        |

One honest note to end on: everything described under V3 and V4 is a plan, not a promise, and none of it is built yet. What is built and tested is V2 — the part nobody sees.